Project Scipio: An Architectural Blueprint for an AI-Powered Battlefield Tactician

Introduction: Architecting the Modern Centurion

In the contemporary and future operational environment, characterized by the rapid proliferation of sensors, the exponential growth of data, and the compression of decision timelines, the traditional paradigms of command and control are increasingly strained. The emerging era of "intelligentized" warfare posits that the ultimate high ground is not geographic, but cognitive. Victory will belong not necessarily to the side with the most firepower, but to the one that can achieve and sustain decision superiority—the ability to sense, understand, decide, and act more quickly and effectively than any adversary. It is within this strategic context that Project Scipio is conceived.

This document provides a comprehensive architectural blueprint for Project Scipio, an advanced artificial intelligence (AI) system designed to function as a grand battlefield tactician. It is crucial to establish from the outset that Scipio is not envisioned as an autonomous commander or a replacement for human leadership. Rather, it is architected as a cognitive enhancement tool, a digital Centurion, designed to augment and empower the human command element. Its purpose is to master the overwhelming volume, velocity, and variety of battlefield information, transforming this chaotic flood of data into actionable intelligence, coherent operational plans, and a decisive tactical edge. The system is intended to be a modular and adaptable director of operations, capable of being modified to fit into any workspace or command environment.

The central thesis of this report is that the creation of a true AI tactician requires a tripartite architecture, a synergistic fusion of three foundational pillars: a **Doctrinal Soul**, a **Data Ecosystem**, and an **AI Core**.

- The **Doctrinal Soul** provides the "why"—the governing principles, logical frameworks, and proven methodologies of warfare that ensure the AI's reasoning is tactically sound, relevant, and aligned with centuries of military experience.
- The **Data Ecosystem** provides the "what"—the all-source, multi-domain torrent of information that fuels the AI's perception and comprehension of the battlespace, forming the basis of its intelligence.
- The **AI Core** provides the "how"—the sophisticated algorithms and computational models that enable the AI to process data, predict outcomes, generate options, and learn from engagement.

This architecture is designed to embody the essential duality of military tactics: the seamless integration of the "art" of creative maneuver and the "science" of technical application. By structuring the system upon this foundation, Project Scipio will provide commanders with a tool that not only calculates odds and optimizes resources but also explores novel strategies and anticipates adversary actions, enabling them to seize and maintain the initiative in the most complex and demanding of environments. This blueprint will detail the theoretical underpinnings, the data requirements, the algorithmic engine, and the operational integration necessary to bring this vision to fruition.

Part I: The Doctrinal Soul of the Machine

The efficacy and trustworthiness of any military AI system are contingent upon its grounding in established military doctrine. An AI that operates outside the logical and ethical frameworks of warfare is not a tool but a liability. Therefore, the foundational layer of Project Scipio—its "Doctrinal Soul"—is an explicit encoding of the principles, processes, and philosophies that govern modern tactical and operational planning. This ensures that the system's outputs are not merely computationally optimal but are also militarily coherent, understandable to its human users, and aligned with the fundamental nature of combat. This section outlines the core tactical functions and decision-making frameworks that will be embedded into the system's architecture, forming the immutable logic upon which all its higher-level cognitive functions will be built.

The Anatomy of Tactical Mastery

To function as a tactician, Project Scipio must first possess a deep, structural understanding of what tactics are. This involves more than recognizing units on a map; it requires a conceptual grasp of the fundamental elements of combat and the principles that guide their application. The system's architecture will be designed to orchestrate the core functions of battle, balance the art and science of their application, and account for the indispensable human dimension of warfare.

Core Functions and Principles

At its most fundamental level, military tactics is the art and science of organizing and employing fighting forces on or near the battlefield. The AI's primary purpose is to design and analyze strategic plans by orchestrating the four interdependent battlefield functions that have defined conflict throughout history: kinetic/firepower, mobility, protection/security, and shock action. The system's design will be architected around the effective synthesis of these functions.

- **Kinetic/Firepower:** This function encompasses the application of destructive force, from direct-fire weapons to long-range artillery and close air support. The AI must possess a comprehensive model of the firepower capabilities of all friendly and enemy systems, including range, lethality, and rate of fire. Its planning modules will use firepower not just for attrition, but also to fix an enemy in place, suppress defenses, and enable decisive maneuver.
- **Mobility:** This function relates to the movement of forces to positions of advantage. The AI will analyze terrain, infrastructure, and enemy dispositions to plan routes for deployment, flanking maneuvers, and rapid exploitation of breakthroughs. It will model the mobility characteristics of every unit type, understanding that achieving a decisive advantage often hinges on being faster than the opponent.
- **Protection and Security:** This function involves the preservation of the force. The AI will integrate measures for cover and concealment, establish security perimeters, and plan for the defense of key assets. It must continuously assess vulnerabilities and recommend actions to mitigate threats, ensuring the force can sustain its combat power.
- **Shock Action:** This is the physical and psychological function of overwhelming an enemy at a decisive point. The AI will analyze opportunities to generate shock through the concentrated application of force, such as a massed armored charge or a surprise infantry

assault. It will recognize that the psychological impact of shock, enhanced by surprise, can be as potent as its physical effects.

The system's effectiveness will be measured by its ability to synchronize these functions into cohesive, combined arms tactics. A key design principle is that for maximum potential to be achieved, all elements of a combined arms team require commensurate levels of mobility, firepower, and protection. The AI will constantly evaluate this balance when composing task forces and planning operations. Furthermore, all AI-generated plans will be evaluated against the foundational Principles of War and Tactical Tenets, such as Mass, Objective, Offensive, Surprise, Economy of Force, Maneuver, Unity of Command, Security, and Simplicity. These principles will serve as a set of high-level heuristics that guide the AI's planning processes, ensuring its recommendations adhere to proven tenets of successful military operations, such as achieving a decision, gaining an advantage, adapting, cooperating, and exploiting success.

The Art and Science of Tactics

A critical insight from modern military doctrine is the division of tactics into two distinct but complementary domains: the "art" and the "science". This doctrinal distinction is not merely philosophical; it provides a direct and powerful architectural pattern for the AI Core of Project Scipio. It dictates the necessity of a hybrid AI system, where different types of AI models are tailored to the unique demands of each domain.

- **The Science of Tactics:** This domain concerns the technical application of combat power. It is the realm of quantifiable data, physical laws, and established procedures. It includes mastering the techniques of marksmanship, gunnery, land navigation, and the precise execution of tasks like calling for close air support. For the AI, this translates to modules that excel at processing and applying hard, empirical data. These modules will be built using predictive analytics, optimization algorithms, and high-fidelity simulations. They will calculate logistics consumption rates, model weapon effectiveness based on range and atmospheric conditions, determine optimal road march speeds, and compute line-of-sight for fields of fire. This is the system's analytical engine, dedicated to achieving "brilliance in the basics" by ensuring every technical detail of a plan is sound and optimized.
- **The Art of Tactics:** This domain involves the creative formation, positioning, and maneuver of combat power in a given situation. It is about answering questions like *when* and *how* to attack—frontally, by flanking, or through envelopment. This creativity is not random; it is a developed capacity acquired through education, practice, and experience. For the AI, this requires models capable of generating novel yet plausible solutions, of identifying an unexpected vulnerability, or of creating the conditions for surprise. This is less about optimizing known variables and more about generating new possibilities. This domain will be powered by generative AI models, such as Large Language Models (LLMs), trained on a vast corpus of military history, doctrine, and past operational plans. These models excel at recognizing deep patterns and combining concepts in new ways, allowing the AI to propose innovative schemes of maneuver that a human planner, constrained by cognitive biases or limited experience, might not consider.

This dual-architecture approach is a core design principle. The "Science" modules provide the rigorous, data-driven validation for the creative proposals generated by the "Art" modules. A plan conceived by the generative "Art" engine will be passed to the analytical "Science" engine for wargaming and refinement, ensuring that a creatively bold concept is also technically feasible and logistically sustainable.

The Human Dimension

A recurring theme in the study of strategy is the recognition that war is a fundamentally human endeavor. A superior tactician understands they are not merely moving chess pieces but are leading real organizations with unique possibilities and constraints. Therefore, Project Scipio must be designed to model and incorporate the "human dimension" into its calculations. An AI that only considers equipment capabilities and terrain will produce brittle, unrealistic plans. The system will include models that account for factors such as:

- **Troop Fatigue and Morale:** Using data from readiness reports and historical performance under similar conditions, the AI will estimate the effects of sustained operations on unit effectiveness. A plan that requires a unit to conduct a forced march after 72 hours of continuous combat will be flagged as high-risk.
- **Training Proficiency:** The system will use the Training (T-level) data from readiness reporting systems to weigh the combat effectiveness of units. A less-equipped but highly trained unit may be assessed as more capable for a complex task than a better-equipped but novice unit.
- **Leadership and Motivation:** The AI will comprehend that a key role of a strategist is to understand people and know how to motivate them. While the AI cannot motivate troops itself, it can provide the human commander with the information needed to do so, such as identifying the critical point in a battle where a commander's presence would have the most impact.
- **The Stresses of War:** The system will be programmed with an understanding of how human beings react to the stresses, agonies, and horrors of war, not just at the tactical level but at the highest levels of command. This will inform its risk assessments, recognizing that plans which place extreme psychological stress on personnel may have a lower probability of success, regardless of the material calculus.

By incorporating these human factors, Project Scipio moves from a simple operations planner to a true strategist's assistant, capable of generating plans that are not only tactically sound but also grounded in the reality of human performance in combat.

The Cognitive Engine - Frameworks for Decision in Conflict

While the principles of tactics provide the AI with its foundational knowledge, it requires structured cognitive frameworks to apply that knowledge in a dynamic environment. Military doctrine has evolved two primary decision-making models that operate at different cadences: the high-speed, adaptive OODA loop for real-time engagement, and the deliberate, methodical Military Decision Making Process (MDMP) for detailed planning. These are not mutually exclusive frameworks; they represent two distinct but interconnected operational modes for the AI. Project Scipio's architecture must support both running in parallel, creating a system that can react instinctively at the tactical edge while simultaneously conducting deep, reflective planning for future operations.

The OODA Loop: The System's Heartbeat

The core real-time processing cycle of Project Scipio will be architected around Colonel John Boyd's Observe, Orient, Decide, Act (OODA) loop. A critical design choice is to model the AI's cognitive engine on Boyd's original, nuanced diagram, not the overly simplified, linear version often depicted. Boyd's complete model is not a simple circle but a complex, multi-layered process with numerous interacting feedback loops. This sophisticated model is uniquely suited to a modern, data-centric environment where information arrives non-sequentially from a multitude of sources in real-time.

The "Orient" phase is the gravitational center of Boyd's model and the most important step in the AI's process. It is here that raw, incoming data from the "Observe" phase is fused with the system's pre-existing knowledge—its understanding of cultural traditions, genetic heritage (doctrine), new information, and previous experiences (analysis of historical battles)—to form a coherent mental model of the current situation. Boyd's diagram explicitly shows that the "Orient" phase is continuously updated by implicit guidance and control from the "Decide" and "Act" phases, as well as by unfolding circumstances from the "Observe" phase. This creates a dynamic, non-linear process of continuous learning and adaptation, which is precisely what is required on a fluid battlefield.

The AI's primary operational goal will be to achieve "decision-action" cycles that are consistently faster, more accurate, and more effective than the adversary's. By processing its own OODA loop more rapidly, the AI can "get inside" the opponent's decision cycle, presenting them with a new reality before they have finished reacting to the old one. This can lead to decision paralysis in the enemy, allowing friendly forces to seize and maintain the initiative—the very definition of an offensive mindset. The OODA loop will therefore function as the system's high-frequency "heartbeat," driving its continuous cycle of situational assessment and tactical adaptation.

The Military Decision Making Process (MDMP): The Deliberate Mind

For more complex, non-instantaneous operational planning, Project Scipio will be structured to support and accelerate the 7-step Military Decision Making Process (MDMP). The MDMP is a methodical, analytical process that provides a proven framework for understanding a problem, developing and comparing solutions, and producing a detailed, executable order. The AI's modules will be explicitly mapped to this doctrinal workflow, acting as a powerful force multiplier for the human planning staff at every stage.

1. **Receipt of Mission:** Upon receiving a mission from a higher headquarters, the AI's Natural Language Processing (NLP) module will parse the order, automatically identifying and cataloging specified and implied tasks, constraints, and the initial commander's intent. It will simultaneously begin a preliminary time analysis, highlighting key deadlines and allocating planning time for the subsequent steps.
2. **Mission Analysis:** This is the most critical step, and where the AI provides the greatest initial benefit. The system will fuse all available intelligence from the data ecosystem to frame the operational environment and the problem. It will identify critical facts and develop necessary planning assumptions, clearly distinguishing between the two. AI-powered analysis will evaluate friendly force capabilities, enemy composition and disposition, and the effects of terrain and weather, providing the staff with a comprehensive foundation for planning. The output of this phase is a proposed problem statement and initial commander's intent for human review and approval.
3. **Course of Action (COA) Development:** Leveraging its generative "Art" module, the AI will develop multiple distinct COAs. Each proposed COA will be internally validated against the five doctrinal criteria: *suitable* (accomplishes the mission), *feasible* (can be done with available resources), *acceptable* (balances cost and risk), *distinguishable* (is significantly different from other COAs), and *complete* (addresses all major tasks). This

allows the planning staff to explore a wider range of options than might be possible under time pressure.

4. **COA Analysis (Wargaming):** Each AI-generated COA undergoes a rigorous wargaming process within the system's high-fidelity simulation environment. The AI will play the friendly COA against the predicted Enemy's Most Likely Course of Action (EMLCOA) and other potential enemy COAs. This process identifies strengths, weaknesses, potential friction points, and likely outcomes for each plan, allowing for refinement and the development of contingencies.
5. **COA Comparison:** The AI will conduct a structured comparison of the wargamed COAs. It will use a decision matrix to score each COA against the commander's evaluation criteria, providing a quantitative analysis of their respective merits and limitations. This data-supported assessment allows the commander to clearly see the trade-offs between different approaches.
6. **COA Approval:** The system will present the recommended COA, along with the second and third-best options, to the commander for a final decision. The presentation will use advanced visualization and explainable AI (XAI) techniques to clearly articulate the rationale behind the recommendation, the key data points that influenced the outcome, and the associated risks. The human commander makes the final, authoritative approval.
7. **Orders Production:** Once a COA is approved, the AI's Natural Language Generation (NLG) module translates the structured plan data into a complete, doctrinally formatted Operation Order (OPORD), including all necessary annexes and supporting graphics. This dramatically accelerates the orders dissemination process, ensuring subordinate units receive clear, actionable instructions with minimal delay.

The system's architecture will treat the OODA loop and the MDMP as parallel, interacting processes. The high-speed OODA cycle provides the real-time inputs for tactical adaptation. A significant change detected during the "Observe" phase of the OODA loop—such as the unexpected appearance of an enemy reserve force—can trigger a "reframing" of the problem within the ongoing MDMP cycle, forcing a reassessment of assumptions and potentially invalidating a previously favored COA. This feedback loop between the system's reflexive and deliberate minds ensures that long-range planning remains continuously grounded in the evolving reality of the battlefield.

Part II: The Data Ecosystem - The Fuel for Cognitive Dominance

An AI is only as intelligent as the data it can access and comprehend. For Project Scipio to function as a grand tactician, it requires a nervous system that extends across the entire battlespace, continuously ingesting, fusing, and structuring information from a vast array of disparate sources. This data ecosystem is the fuel for the system's cognitive engine; its quality, timeliness, and comprehensiveness directly determine the quality of the AI's perception, comprehension, and prediction. This section details the architecture for this ecosystem, defining the data sources, structures, and models necessary to create a rich, multi-dimensional, and machine-readable representation of the operational environment.

The Sentient Battlefield - Achieving Pervasive Situational Awareness

The foundational output of the data ecosystem is Situational Awareness (SA)—the ability to observe and anticipate activities in the area of interest, or more simply, to know what is happening on the battlefield. Project Scipio's SA module will be architected around the formal three-level model of SA, a hierarchical framework that ensures the system progresses from raw data points to genuine predictive insight. This structured approach is critical for transforming a flood of information into actionable understanding.

- **Level 1: Perception of the Elements:** This is the most basic level of SA, involving the detection, recognition, and monitoring of key elements within the time and space of the operational environment. The AI's perception layer is responsible for ingesting and processing raw data from all available sources to identify and track the current state of these elements: the locations, conditions, and actions of friendly forces, enemy forces, neutral entities, and significant environmental factors. This is a task of massive data fusion, integrating real-time streams from Intelligence, Surveillance, Target Acquisition, and Reconnaissance (ISTAR) platforms, signals intelligence (SIGINT), open-source intelligence (OSINT), and friendly force tracking systems to build a foundational, dynamic map of the battlespace.
- **Level 2: Comprehension of the Situation:** Moving beyond simple perception, comprehension involves synthesizing the data from Level 1 to understand its meaning, significance, and relevance in the context of the mission and the commander's goals. The AI's comprehension module analyzes the relationships between the perceived elements. It does not just see a friendly tank platoon and an enemy anti-tank position; it comprehends that the enemy position is a direct threat to the friendly unit's avenue of approach. This level involves identifying patterns, assessing the current situation against the planned operation, and understanding the "why" behind the "what". It is at this level that raw data is transformed into tactically relevant information.
- **Level 3: Projection of Future Status:** This is the highest and most valuable level of SA, where the system uses the comprehensive understanding developed in Level 2 to anticipate the future state of the operational environment. The AI's projection module is its predictive core. By extrapolating the dynamics of the current situation forward in time, it can forecast likely enemy movements, predict potential threats before they materialize, and identify fleeting opportunities for exploitation. Achieving this level of SA is the key to seizing the initiative from the adversary. It allows the commander to act proactively rather than reactively, shaping the future of the battle instead of merely responding to it. The ability to "get inside" the opponent's OODA loop is fundamentally dependent on achieving and maintaining Level 3 SA.

The All-Source Intelligence Fabric

To populate the three levels of Situational Awareness, Project Scipio requires an intelligence fabric capable of weaving together data streams from every conceivable source into a single, cohesive tapestry. The architecture must be designed for extreme heterogeneity, capable of ingesting everything from high-frequency sensor data and structured logistics reports to unstructured doctrinal texts and public social media feeds.

Real-Time Sensor and Intelligence Data

The system's perception of the current battlespace is built upon a continuous stream of data from a multi-domain sensor network.

- **ISTAR Data:** The system will fuse data from a wide array of platforms, including satellites, uncrewed aircraft systems (UAS), and ground sensors. This requires the ability to process multiple sensor modalities, each providing a unique layer of information: high-resolution optical and infrared (EO/IR) video for visual identification; Synthetic Aperture Radar (SAR) and Moving Target Indicator (MTI) data for detecting vehicles through weather and foliage; and Signals Intelligence (SIGINT) to detect and locate enemy communications and radar emitters. A key capability will be the real-time processing of full-motion video (FMV) from drone feeds to enable Automatic Target Recognition (ATR), reducing the cognitive load on human analysts.

- **Geospatial Intelligence (GEOINT):** GEOINT provides the fundamental geospatial context—the digital map—upon which all other intelligence is layered. This is not a static backdrop but a dynamic, multi-layered data set critical for all aspects of planning and execution. Core GEOINT data includes Digital Elevation Models (DEMs) for line-of-sight analysis and route planning; detailed data on vegetation, hydrology, and soil composition that affect mobility; and a comprehensive library of man-made infrastructure such as roads, bridges, and buildings. Crucially, this also includes real-time and forecast weather data—temperature, precipitation, cloud cover, visibility, and wind—as environmental conditions have a profound impact on both military operations and the performance of sensors like EO/IR systems.
- **Open-Source Intelligence (OSINT):** In modern conflicts, the information environment is a critical domain. The AI will incorporate an OSINT module to harvest and analyze publicly available information. This includes monitoring social media platforms (SOCMINT) to gauge civilian sentiment and track disinformation campaigns, analyzing commercial satellite imagery to monitor infrastructure and activity in non-hostile areas, and scanning news media and other sources for information that provides context to the operational environment.

Friendly Force and Logistics Data

A complete operational picture requires an equally detailed and timely understanding of the status of one's own forces. An AI cannot develop a feasible plan without knowing the precise capabilities and limitations of the tools at its disposal.

- **Force Strength and Disposition:** The system will maintain a continuous, real-time plot of all friendly assets by integrating data from Blue Force Tracking (BFT) systems and Identification Friend or Foe (IFF) networks. This common operational picture is essential for command and control, coordination of fires and maneuver, and, critically, the prevention of fratricide.
- **Readiness and Logistics Status:** A unit's presence on the map does not equate to its actual combat capability. Project Scipio must integrate with authoritative military data systems to get a true assessment of readiness. This includes pulling data from readiness reporting systems like the Defense Readiness Reporting System (DRRS) and logistics information systems (LIS) such as the Standard Army Ammunition System (SAAS), Transportation Coordinators' – Automated Information for Movements System II (TC-AIMS II), and Property Book Unit Supply Enhanced (PBUSE). This provides the AI with granular data on the four critical resource areas: Personnel (P-level), Supply/Equipment on Hand (S-level), Equipment Condition/Readiness (R-level), and Training (T-level). This data is what allows the AI to assess the *actual* combat power of a unit, enabling it to generate plans that are not just theoretically sound but practically achievable.

Doctrinal and Historical Data

The AI's "knowledge base"—the foundation of its Orient phase and its generative capabilities—must be built from a vast corpus of textual and historical data.

- **Military Doctrine:** A comprehensive library of digitized military doctrine, including field manuals, joint publications, and Tactics, Techniques, and Procedures (TTPs), will be used to train the AI's core models. This data provides the fundamental rules, principles, and vocabulary of warfare, ensuring the AI "thinks" and "speaks" in a way that is doctrinally sound and understandable to military professionals.
- **Historical Battle Data:** The system will be trained on a massive dataset of historical conflicts. This data, curated from archives and after-action reviews, will include detailed information on unit dispositions, orders issued, actions taken, and the resulting outcomes. By analyzing thousands of past engagements, the AI's machine learning models can identify deep patterns in adversary behavior, learn the relative effectiveness of different tactics under specific conditions, and build the predictive models necessary to anticipate an enemy's most likely course of action.

The following matrix provides a structured overview of the primary data sources required by Project Scipio, detailing their format, frequency, and relevance to the system's core modules. This serves as an actionable specification for the data architects and engineers responsible for constructing the system's data ingestion and integration pipelines.

Battlefield Data Source Matrix

Data Category	Specific Data Type	Example Source Systems	Data Format / Standard	Update Frequency	Relevance to AI Modules
ISTAR	Electro-Optical/ Infrared (EO/IR) Video	UAS (Predator, Reaper), Satellites	Full Motion Video (FMV) streams, NITF	Real-Time	SA (L1), Automatic Target Recognition
	Signals Intelligence (SIGINT)	SIGINT platforms, sensors	Raw I/Q data, specific formats	Real-Time	SA (L1, L2), Enemy Comms Analysis
	Radar (SAR, MTI)	JSTARS, Global Hawk	Standardized radar data formats	Near Real-Time	SA (L1, L2), Enemy Movement Detection
GEOINT	Terrain & Elevation	NGA, Commercial Providers	Digital Elevation Model (DEM), DTED	Static/Periodic	SA (L2, L3), COA Dev, Maneuver Planning
	Weather	Military Weather Services	METAR, GRIB2	Real-Time/Forecast	SA (L2, L3), COA Analysis (effects on sensors/ops)
	Hydrography/In	AGO, NGA,	Vector data	Static/Periodic	SA (L2),

Data Category	Specific Data Type	Example Source Systems	Data Format / Standard	Update Frequency	Relevance to AI Modules
	frastructure	OpenStreetMap	(Shapefiles), CMAP		Obstacle Analysis, Logistics Planning
Force Status	Friendly Force Tracking	BFT, JBC-P	TDL (e.g., Link 16 J-Series)	Real-Time	SA (L1, L2), Fratricide Prevention, C2
	Enemy Order of Battle (OOB)	Intelligence Fusion Centers	SALUTE/DRAWD reports, structured data	Periodic/Event-driven	SA (L2, L3), Enemy COA Prediction, COF Calc
	Readiness (P, S, R, T)	DRRS, GCSS	XML, Database records	Daily/Weekly	COA Feasibility, Resource Management
Logistics	Supply & Ammunition Status	LIS (SAAS, PBUSE)	Database records, custom formats	Near Real-Time	Resource Management, COA Feasibility
	Transportation & Movement	TC-AIMS II	RFID data, satellite tracking	Near Real-Time	Logistics Planning, Sustainment Analysis
Knowledge Base	Military Doctrine	Digital Libraries (e.g., AFDP)	PDF, XML, structured text	Static	COA Generation, Wargaming Ruleset
	Historical Battle Data	Archives, After-Action Reviews	Structured event logs, text	Static	Predictive Model Training (Enemy Intent)

The Unified Battlespace Model

To effectively manage and exploit this diverse data ecosystem, Project Scipio will depart from traditional, siloed database architectures. Instead, its core data structure will be a **Unified Battlespace Model**, implemented as a dynamic, in-memory graph database. In this model, every entity in the operational environment—a friendly tank platoon, an enemy radar emplacement, a bridge, a weather front, a doctrinal principle—is represented as an object or "node." These nodes possess a rich set of attributes (e.g., location, speed, status, capabilities, fuel level, ammunition count).

Crucially, the power of this model lies in the "edges" that define the relationships between these nodes (e.g., "is targeting," "is providing support to," "is constrained by," "is communicating with"). This graph structure is highly flexible and is natively optimized for the types of complex queries and deep pattern analysis required by the AI Core. For example, the system could execute a query to find "all enemy air defense systems (nodes) that are within range of (relationship) the primary air corridor (node) and are not currently being targeted by (relationship) any friendly

SEAD assets (nodes)." This type of relational query is computationally expensive in traditional databases but is highly efficient in a graph model. This unified model provides a single, coherent, and continuously updated digital twin of the battlespace, serving as the "single source of truth" for all of the AI's cognitive processes.

Part III: The AI Core - Forging the Tactical Intellect

The AI Core is the cognitive engine of Project Scipio, where the fused data from the ecosystem is transformed into understanding, prediction, and actionable plans. This is the heart of the system, comprising a suite of specialized AI and machine learning (ML) models designed to replicate and accelerate the reasoning processes of an expert military staff. The architecture of the AI Core is explicitly mapped to the OODA loop and MDMP frameworks, with distinct modules dedicated to the functions of sense-making (Orient), decision optimization (Decide), and order generation (Act). This section details the specific AI methodologies that will power these modules, creating a system capable of both analytical rigor and creative insight.

The "Orient" Module - AI-Powered Sense-Making and Prediction

The "Orient" phase is the most critical component of the OODA loop, as it shapes how an entity perceives and interacts with its environment. The AI's Orient module is designed to be a continuous sense-making engine, constantly working to transform the torrent of data from the Unified Battlespace Model into a coherent and predictive understanding of the operational situation. This goes far beyond simply plotting icons on a map; it is about uncovering the underlying dynamics of the battlefield and anticipating what will happen next.

Pattern Recognition and Anomaly Detection

The first task of the Orient module is to establish a baseline of normalcy for the operational environment. Using unsupervised learning techniques such as clustering algorithms and autoencoders, the AI will analyze vast amounts of sensor and intelligence data to identify recurring "patterns of life". This could include normal traffic patterns on key roads, typical communication frequencies and volumes from enemy units, or standard civilian activity in urban areas. Once these baselines are established, the system can apply anomaly detection algorithms to flag any significant deviations. An unexpected convergence of vehicles, a sudden silence from a normally active command net, or the absence of civilian traffic in a bustling market could all be early indicators of an impending threat or a developing opportunity. By automatically surfacing these anomalies, the AI focuses the attention of human analysts on the most critical and time-sensitive information, preventing vital clues from being lost in the noise.

Predictive Analysis and Intent Recognition

The ultimate goal of the Orient phase is to move from reaction to anticipation. The most powerful function of this module is its ability to predict the Enemy's Most Likely Course of Action (EMLCOA). This is achieved through a suite of supervised learning models, including deep neural networks and Bayesian inference networks, trained on the extensive historical battle data and digitized enemy doctrine stored in the system's knowledge base.

The process begins by analyzing the current enemy disposition, activity, and composition (the

"SALUTE" report). The AI then compares this current state to thousands of historical examples and doctrinal templates. For instance, if the AI observes an enemy armored formation adopting a specific wedge formation, supported by self-propelled artillery and forward-deployed reconnaissance assets, it can reference its training data to infer that this is a doctrinal precursor to a deliberate breach operation. The system would then predict the likely objective, axis of advance, and timing of the enemy attack. This predictive capability allows friendly forces to shift from a reactive posture to a proactive one, enabling them to reinforce the threatened sector, prepare counter-attacks, and shape the engagement on their own terms, thereby seizing the initiative.

The "Decide" Module - Algorithmic Wargaming and COA Optimization

Once the Orient module has developed a rich, predictive understanding of the situation, the Decide module is responsible for generating and evaluating friendly Courses of Action (COAs) to achieve the mission's objectives. This module employs a novel, hybrid AI approach that directly mirrors the doctrinal division between the "art" and "science" of tactics, ensuring that the COAs it produces are both creatively conceived and rigorously tested.

Hybrid COA Development

The system utilizes a two-stage process for COA development, leveraging the distinct strengths of different AI paradigms.

- **Stage 1: Generative AI for Creative COA Generation:** The process begins with a Large Language Model (LLM) that has been extensively fine-tuned on a specialized corpus of military doctrine, TTPs, historical battle plans, and after-action reviews. When given the commander's intent, the mission statement, and the output of the Orient module (the current SA and EMLCOA), this generative model will produce a diverse set of tactically plausible COAs. It will output these COAs in a structured format, including a narrative concept of the operation, a task organization, and a schematic sketch of the maneuver. The key advantage of using generative AI here is its ability to rapidly explore a vast solution space and propose novel combinations of tactical actions, helping to overcome the cognitive biases and "groupthink" that can sometimes constrain human planning teams under pressure.
- **Stage 2: Reinforcement Learning for Rigorous COA Analysis:** Each COA generated in Stage 1 is then subjected to a rigorous wargaming process. The COA serves as the high-level policy guiding a Reinforcement Learning (RL) agent within a high-fidelity, physics-based simulation of the battlespace. The RL agent executes the plan against the simulated enemy (which is operating based on the EMLCOA), exploring thousands of variations and tactical encounters at machine speed. The simulation's reward function is multi-objective, designed to train the agent to optimize for mission accomplishment while minimizing friendly casualties and logistical expenditure.

A critical component of this reward function is a digital, real-time **Correlation of Forces (COF) Calculator**. Based on doctrinally accepted methodologies, this calculator provides a quantitative score of the relative combat power between friendly and enemy forces at any given point in the simulation, taking into account unit type, equipment, and current strength. Engagements that result in favorable force ratios receive a positive reward, while those that lead to unfavorable or attritive conditions are penalized. This ensures the RL agent learns to generate tactical actions that are not just successful, but efficient and consistent with sound military principles. This

wargaming process rigorously tests each creative COA, pruning suboptimal variations and identifying the plan with the highest overall probability of success.
The following table explicitly links these and other AI techniques to each step of the MDMP, providing a clear roadmap for how the AI Core supports the entire planning cycle.

AI Model Application in the MDMP

MDMP Step	Primary AI Technique(s)	Function / Purpose	Supporting Snippets
1. Receipt of Mission	Natural Language Processing (NLP)	Parse OPORDs, extract tasks, constraints, and commander's intent.	
2. Mission Analysis	Unsupervised & Supervised Learning	Cluster intelligence data, identify patterns, predict EMLCOA, frame the problem.	
3. COA Development	Generative AI (Fine-tuned LLM)	Generate multiple, distinct, and doctrinally sound COA narratives and sketches.	
4. COA Analysis	Reinforcement Learning (RL), Simulation	Wargame each generated COA in a simulated environment to test for robustness and outcomes.	
5. COA Comparison	Predictive Analytics, COF Calculation	Quantitatively score and rank wargamed COAs based on success probability, cost, and risk.	
6. COA Approval	Explainable AI (XAI), Visualization	Present recommended COA and its rationale to the commander in an intuitive, understandable format.	
7. Orders Production	Natural Language Generation (NLG)	Automatically draft a complete, formatted OPORD from the approved COA data structure.	

The "Act" Module - From Optimal Decision to Actionable Orders

The final stage of the AI Core's process is to translate the commander's decision into action. Once the human commander reviews the results of the wargaming and comparison and provides their final approval for a COA, the "Act" module takes over. This module utilizes Natural Language Generation (NLG) techniques, a counterpart to the NLP used in mission receipt. The system takes the highly structured data that defines the approved COA—the task organization, unit-specific tasks, control measures, timelines, and logistical plan—and automatically generates a complete, human-readable, and doctrinally correct Operation Order (OPORD). This process, which can take human staffs hours to complete, is accomplished in seconds. This capability drastically reduces the time between decision and execution, a critical factor in gaining and maintaining a tempo advantage over the adversary. The generated order is then disseminated to subordinate units over the appropriate tactical communication networks.

Part IV: The Operational Architecture - System Resilience and Human Command

A theoretically brilliant AI is of no value if it cannot function amidst the chaos, friction, and uncertainty of a real-world battlefield. The operational architecture of Project Scipio addresses the practical challenges of implementation, focusing on its seamless integration with existing military systems, its resilience in the face of a contested and degraded environment, and, most importantly, the design of a command interface that ensures the human commander remains firmly in control. This section details the hardware, networking, and human-computer interaction principles that will make Project Scipio a robust, reliable, and trustworthy tool of command.

C4ISR Integration and Network Architecture

Project Scipio is not a standalone system; it is a component within a larger Command, Control, Communications, Computers, Intelligence, Surveillance, and Reconnaissance (C4ISR) ecosystem. Its ability to function depends on its capacity to seamlessly exchange data with a wide variety of joint and coalition platforms.

Tactical Data Links (TDL)

The system's primary means of communication with other military assets will be through standard Tactical Data Links (TDLs). The architecture will include native support for Link 16, the NATO standard for jam-resistant, high-capacity data exchange used by most modern air, sea, and ground platforms. By communicating via Link 16's J-series message formats, Project Scipio can receive real-time friendly and enemy track data and disseminate control messages and situational awareness updates across the network. The system will also support IP-based protocols like the Standard Interface for Multiple Platform Link Evaluation (SIMPLE), which enables the transmission of TDL messages over modern digital networks, ensuring forward compatibility.

Secure Communications

Given the sensitive nature of command and control data, all communications to, from, and within the Project Scipio system will be protected by robust, multi-layered security protocols. All data transmissions will utilize end-to-end encryption, with AES-256 as the minimum standard, to prevent interception and manipulation by adversaries. To operate in a contested electromagnetic spectrum, the system's radio communications will employ advanced electronic protection measures, including frequency-hopping spread spectrum (FHSS) and low-probability-of-intercept/low-probability-of-detection (LPI/LPD) waveforms, making them highly resilient to enemy jamming and eavesdropping efforts.

Resilience by Design - Mastering the Contested Environment

The modern battlefield is an inherently hostile environment for digital systems. Adversaries will actively seek to disrupt communications through electronic warfare, and the physical realities of combat operations often result in unreliable or non-existent network connectivity. The architecture of Project Scipio must therefore be fundamentally resilient, designed from the ground up to operate effectively in a Disconnected, Intermittent, and Low-Bandwidth (DIL) environment.

Operating in DIL Environments

To counter the threat of connectivity loss, Project Scipio will be built on a decentralized, distributed architecture rather than a centralized, cloud-dependent one. While a rear-echelon command post may host the full suite of AI models, forward-deployed command nodes will be equipped with powerful "edge computing" hardware capable of running mission-critical SA and planning modules locally. This ensures that core functionality remains available even if the link to higher headquarters is severed.
The system will employ several strategies to manage data in these constrained environments:

- **Data Prioritization:** Using Quality of Service (QoS) protocols, the system will prioritize the transmission of the most critical data streams. For example, an urgent enemy contact report or a call for fire would take precedence over a routine logistics update, ensuring that essential information gets through even on a low-bandwidth link.
- **Asynchronous Replication:** The system will use asynchronous data replication and principles from Delay/Disruption Tolerant Networking (DTN). Instead of requiring a constant connection, data packets can be stored locally and forwarded to other nodes when a connection becomes available (a "store-and-forward" model). This allows the distributed network to maintain a coherent operational picture over time, even with only sporadic connectivity between nodes.

Fault Tolerance and Redundancy

The system must be designed to withstand the loss of individual components without catastrophic failure. The principle of fault tolerance will be embedded throughout the architecture, using both hardware and software redundancy.

- **Space Redundancy:** Critical hardware components, such as servers and network routers, will have physical backups that can be automatically brought online in the event of a failure.
- **Software and Data Redundancy:** Key C2 functions and the Unified Battlespace Model will be distributed and replicated across multiple nodes in the network. If a primary command node is destroyed, a secondary node can seamlessly take over its functions, ensuring continuity of operations. This approach allows for "graceful degradation"—where the loss of a component might reduce the system's overall capacity but does not cause a total system failure—rather than creating a single point of failure that could cripple the entire force.

The Command Interface - Ensuring Meaningful Human Control

The most sophisticated AI is useless if its human user cannot understand, trust, and effectively control it. The use of AI in decisions that may lead to the use of lethal force carries profound ethical weight, and the design of Project Scipio must be guided by this reality. The greatest non-technical risk to the system is not the specter of rogue AI, but the far more insidious danger of human abdication of responsibility through automation bias. Therefore, the command interface is not a secondary feature; it is a primary control mechanism for the ethical and effective employment of the entire system.

Ethical Imperatives and Meaningful Human Control (MHC)

The unassailable principle guiding the system's design is that of **Meaningful Human Control (MHC)**. Project Scipio is a decision *support* tool, not a decision *maker*. The human commander always retains ultimate authority, responsibility, and moral accountability for every action taken. The system must be designed to enhance the commander's judgment, not supplant it. Every aspect of the user interface and workflow must be engineered to reinforce this principle.

Human-in-the-Loop (HITL) Design

The architecture will be explicitly Human-in-the-Loop (HITL) at every critical decision point. The system is not permitted to autonomously execute lethal actions or issue operational orders. Its function is to generate recommendations and present them to the commander for approval, modification, or rejection. The command interface must provide the tools necessary for the commander to exercise this control effectively. This includes the ability to easily interrogate the AI's reasoning, view the specific data and assumptions that underpin a recommendation, understand the system's confidence level in its own analysis, and seamlessly override any suggestion.

Mitigating Automation Bias

Automation bias is the well-documented human tendency to over-rely on automated systems, often accepting their outputs without sufficient critical scrutiny. A poorly designed interface that simply presents "The Optimal Solution" would actively encourage this dangerous bias. The command interface for Project Scipio must therefore be designed *defensively* to combat this tendency and promote critical engagement.

- **Presenting Options, Not Answers:** Instead of displaying a single "best" COA, the interface will present the commander with the top-ranked options (e.g., the top three). For each option, it will display not just the projected probability of success, but also a detailed breakdown of the associated risks, logistical costs, and key assumptions. This forces the commander to engage in a comparative analysis and apply their own judgment and risk tolerance to the decision.
- **Assumption and Uncertainty Visualization:** The interface will feature an "Assumption Monitor" that explicitly lists the critical assumptions the AI has made in its analysis. For example, it might state: "This plan assumes the enemy's reserve is located at Grid X with 85% confidence. If this assumption is false, the plan's success probability drops to 40%." This immediately highlights for the commander the most critical Priority Intelligence

Requirements (PIRs) that need to be confirmed, prompting them to direct reconnaissance assets to validate the AI's assumptions before committing forces.
- **Explainable AI (XAI):** The system will incorporate XAI techniques to make the AI's "thought process" as transparent as possible. When the AI recommends a particular flanking maneuver, the commander can query the system, which might respond by visually highlighting the terrain data, enemy disposition, and doctrinal principle (e.g., "Attack an unguarded flank") that led to its conclusion.

This approach transforms the command interface from a passive display into an active tool for a critical dialogue between the human and the AI. It enforces the HITL model by design, making the preservation of Meaningful Human Control an active, engineered feature of the system, ensuring that Project Scipio remains a powerful instrument in the hands of a responsible human commander.

Conclusion and Strategic Recommendations

Project Scipio represents an architectural vision for the future of military command and control—a future where human intellect is augmented, not replaced, by the power of artificial intelligence. This blueprint has detailed a comprehensive, multi-layered system designed to provide commanders with a decisive cognitive edge on the modern battlefield. The strength of this architecture lies not in any single algorithm or data source, but in the synergistic fusion of its three core pillars: a **Doctrinal Soul** that ensures military coherence, a vast **Data Ecosystem** that provides unparalleled situational awareness, and a hybrid **AI Core** that combines analytical rigor with creative insight.

The system's cognitive engine, built upon the dual frameworks of the adaptive OODA loop and the methodical MDMP, allows it to operate effectively across the full spectrum of tactical and operational challenges. Its data-centric design, culminating in a Unified Battlespace Model, transforms the chaos of battlefield information into a structured, machine-readable reality. The novel, hybrid AI approach to Course of Action development—using generative models to create options and reinforcement learning to test them—promises to deliver plans that are not only optimized but also innovative. Finally, the unwavering commitment to a resilient, fault-tolerant operational architecture and a command interface built on the principle of Meaningful Human Control ensures that this powerful technology remains a robust and ethical tool in the hands of the warfighter.

To translate this architectural vision into an operational capability, the following strategic recommendations are proposed:

1. **Adopt a Phased, Modular Development Approach:** The complexity of Project Scipio necessitates an iterative and modular development strategy. The initial phase should focus on establishing the foundational **Data Ecosystem**. This involves building the data ingestion pipelines, developing the Unified Battlespace Model, and perfecting the Level 1 and Level 2 Situational Awareness modules. This initial capability, even without the advanced AI Core, would provide immense value to command staffs by delivering a fused, common operational picture.
2. **Prioritize the Development of the "Orient" Module:** The next phase should focus on the predictive capabilities of the AI Core. Developing the pattern recognition and EMLCOA prediction models will provide commanders with the Level 3 Situational Awareness—the ability to anticipate—that is the key to seizing the initiative.
3. **Invest in Curated, High-Quality Training Data:** The performance of the AI Core is

wholly dependent on the quality of its training data. A significant, dedicated effort must be made to collect, digitize, and structure the vast amounts of military doctrine and historical battle data required to train the generative and predictive models. This is not a secondary task; it is a primary prerequisite for success.
4. **Co-Design the Command Interface with End-Users:** The development of the human-machine interface must be a collaborative process from day one, involving experienced commanders and staff officers. Their insights are essential to creating a system that is intuitive, trustworthy, and actively mitigates automation bias, ensuring that the principles of Meaningful Human Control are woven into the fabric of the system.
5. **Establish a Robust Testing and Validation Framework:** The system must be rigorously tested in high-fidelity simulations and live training exercises against skilled human opponents. This is essential not only for refining the AI models but also for building the trust and confidence of the operators who will one day rely on it in combat.

Project Scipio is an ambitious undertaking, but the strategic imperative is clear. In an era where the speed and quality of decision-making are the most critical determinants of victory, we must equip our leaders with the tools to out-think and out-pace any adversary. By grounding advanced technology in proven doctrine and placing it in the service of human command, Project Scipio provides a viable and necessary path toward ensuring decision superiority on the battlefields of today and tomorrow.

Works cited

1. China's Military Employment of Artificial Intelligence and Its Security Implications, https://www.iar-gwu.org/print-archive/blog-post-title-four-xgtap 2. TACTICAL FUNDAMENTALS B2B2269 STUDENT HANDOUT, https://www.trngcmd.marines.mil/Portals/207/Docs/TBS/B2B2269%20Tactical%20Fundamentals.pdf?ver=2015-06-08-142423-523 3. Military tactics - Wikipedia, https://en.wikipedia.org/wiki/Military_tactics 4. What are military tactician jobs like? - CollegeVine, https://www.collegevine.com/faq/130811/what-are-military-tactician-jobs-like 5. Course of Action Development, Defense, https://irp.fas.org/doddir/army/iobc/coadefso.htm 6. What's the Matter with Being a Strategist? - USAWC Press - Army War College, https://press.armywarcollege.edu/cgi/viewcontent.cgi?article=2552&context=parameters 7. DRRS-MC Commander's Readiness Handbook - Marine Corps University, https://www.usmcu.edu/Portals/218/Commanders_Readiness_Handbook_2020-09-16.pdf 8. www.ebsco.com, https://www.ebsco.com/research-starters/military-history-and-science/ooda-loop#:~:text=The%20OODA%20loop%20is%20a%20decision%2Dmaking%20concept%20that%20was,and%20moving%20to%20take%20action. 9. OODA loop | Research Starters - EBSCO, https://www.ebsco.com/research-starters/military-history-and-science/ooda-loop 10. OODA Loop: A Blueprint for the Evolution of Military Decisions, https://www.rti.com/blog/ooda-loop-a-blueprint-for-the-evolution-of-military-decisions 11. The OODA Loop - The Decision Lab, https://thedecisionlab.com/reference-guide/computer-science/the-ooda-loop 12. The OODA Loop -- Observe, Orient, Decide, Act - LessWrong, https://www.lesswrong.com/posts/hgttKuASB55zjoCKd/the-ooda-loop-observe-orient-decide-act 13. OODA loop - Wikipedia, https://en.wikipedia.org/wiki/OODA_loop 14. Situation awareness - Wikipedia, https://en.wikipedia.org/wiki/Situation_awareness 15. Military Decision-Making Process (MDMP): 7 Steps | Garmont Tactical,

https://garmonttactical.com/post/military-decision-making-process-mdmp-7-steps-from-intel-to-action.html 16. No. 15-06 MDMP MAR 15 - Army University, https://armyuniversity.edu/cgsc/cgss/files/15-06_0.pdf 17. Course of action (COA) - Benefits.com, https://benefits.com/glossary/course-of-action-coa/ 18. This report documents the research on course of action (COA) analysis in alignment with commander's intent - dodccrp.org, http://www.dodccrp.org/events/2006_CCRTS/html/papers/036.pdf 19. AFDP 5-0, Planning - Air Force Doctrine, https://www.doctrine.af.mil/Portals/61/documents/AFDP_5-0/AFDP5-0Planning.pdf 20. flashpoint.io, https://flashpoint.io/blog/what-is-situational-awareness/#:~:text=Military%20situational%20awareness%20is%20the,is%20happening%20on%20the%20battlefield. 21. Situational Awareness in Defense Industry | Definition & Benefits - Getac, https://www.getac.com/en/industries/defense/situational-awareness/ 22. How Military Teams Use OSINT to Boost Threat Intelligence and Awareness - Rishi Sec, https://rishisec.com/how-military-teams-use-osint-to-boost-threat-intelligence-and-battlefield-awareness/ 23. Intelligence, surveillance, target acquisition, and reconnaissance - Wikipedia, https://en.wikipedia.org/wiki/Intelligence,_surveillance,_target_acquisition,_and_reconnaissance 24. STRENGTHENING ISR CAPABILITIES WITH RF SENSOR PAYLOADS IN UNMANNED AERIAL VEHICLES - CRFS, https://www.crfs.com/hubfs/Brochures/ISR%20capabilities%20with%20RF%20sensor%20payloads%20in%20UAVs.pdf 25. What Is Situational Awareness? A Crash Course For Intelligence Teams | Flashpoint, https://flashpoint.io/blog/what-is-situational-awareness/ 26. Enhancing situational awareness in the military - Viasat, https://www.viasat.com/perspectives/government/2024/enhancing-situational-awareness-in-the-military/ 27. Logistics Information Systems (LIS) - PEO Enterprise, https://www.eis.army.mil/programs/lis 28. Machine Intelligence Technology for Automatic Target Recognition - MIT Lincoln Laboratory, https://archive.ll.mit.edu/publications/journal/pdf/vol02_no2/2.2.8.machineintelligence.pdf 29. Deep learning assisted Computer Vision and Visual Analytics in Autonomous Unmanned Aerial Vehicles - IEEE Aerospace and Electronic Systems Society, https://ieee-aess.org/files/ieeeaess/2023-04/Deep%20learning%20assisted%20Computer%20Vision%20and%20Visual%20Analytics%20in%20Autonomous%20Unmanned%20Aerial%20Vehicles%20%282%29%20%281%29%20%281%29.pdf 30. Mission Planning & Terrain Analysis in Complex Environments ..., https://garmonttactical.com/post/mission-planning-and-terrain-analysis-how-modern-us-army-units-navigate-complex-environments.html 31. What Is Geospatial Intelligence and How Has It Evolved? | American Public University, https://www.apu.apus.edu/area-of-study/intelligence/resources/what-is-geospatial-intelligence/ 32. Army Geospatial Center > What we do > Terrain, https://www.agc.army.mil/What-we-do/Terrain/ 33. The Role of Geospatial Intelligence in Modern Military Operations - EarthArXiv, https://eartharxiv.org/repository/view/8618/ 34. Solutions for Governments and Defense | Meteomatics, https://www.meteomatics.com/en/governments-and-defense/ 35. air force doctrine publication (afdp) 3-59 weather operations, https://www.doctrine.af.mil/Portals/61/documents/AFDP_3-59/3-59-AFDP-WEATHER-OPS.pdf 36. Intelligence Discovery Experiment - Line of Departure, https://www.lineofdeparture.army.mil/Journals/Military-Intelligence/MIPB-January-June-2024/Intelligence-Discovery-Experiment/ 37. Identification friend or foe - Wikipedia, https://en.wikipedia.org/wiki/Identification_friend_or_foe 38. GAO- 21-99 Highlights, CLOSE AIR SUPPORT: Actions Needed to Enhance Friendly Force Tracking Capabilities and Fully Evaluate Trai - Government Accountability Office, https://www.gao.gov/assets/gao-21-99-highlights.pdf 39. Order of battle - Wikipedia, https://en.wikipedia.org/wiki/Order_of_battle 40. Defense Readiness Reporting System – Strategic (DRRS-S) - DOT&E, https://www.dote.osd.mil/Portals/97/pub/reports/FY2016/dod/2016drrs.pdf?ver=2019-08-22-.... 41. MCO 3000.11E > United States Marine Corps Flagship > Electronic Library Display, https://www.marines.mil/News/Publications/MCPEL/Electronic-Library-Display/Article/898910/mco-300011e/ 42. The Impact Of Artificial Intelligence On The Military Decision-Making Process And Mission Command - TDHJ.org, https://tdhj.org/blog/post/ai-military-decision-making/ 43. Air Force Doctrine Note 25-1, Artificial Intelligence (AI), https://www.doctrine.af.mil/Portals/61/documents/AFDN_25-1/AFDN%2025-1%20Artificial%20Intelligence.pdf 44. Battlefield Evidence in the Age of Artificial Intelligence-Enabled Warfare, https://cjil.uchicago.edu/print-archive/battlefield-evidence-age-artificial-intelligence-enabled-warfare 45. Data Dominance in Modern Warfare The Crucial Role of AI and Data Analytics - The Forge, https://theforge.defence.gov.au/article/data-dominance-modern-warfare-crucial-role-ai-and-data-analytics 46. How AI with 'nurtured consciousness' could transform warfare ..., https://www.atlanticcouncil.org/blogs/new-atlanticist/how-ai-with-nurtured-consciousness-could-transform-warfare/ 47. The Role of AI and Machine Learning in Military Decision-Making Software | Attract Group, https://attractgroup.com/blog/the-role-of-ai-and-machine-learning-in-military-decision-making-software/ 48. Modernizing Military Decision-Making: Integrating AI into Army Planning, https://www.armyupress.army.mil/Journals/Military-Review/Online-Exclusive/2025-OLE/Modernizing-Military-Decision-Making/ 49. Artificial Intelligence for Decision Support in Command and Control Systems - FOI, https://foi.se/download/18.41db20b3168815026e010/1548412090368/Artificial-intelligence-decision_FOI-S--5904--SE.pdf 50. AI Agents for Course of Action (COA) Generation:, https://s47507.pcdn.co/wp-content/uploads/2025/06/Seekr-Agentic-AI-for-COA-Generation-eBook.pdf 51. Integrating Artificial Intelligence and Machine ... - USAWC Press, https://press.armywarcollege.edu/cgi/viewcontent.cgi?article=1976&context=monographs 52. Automated Course of Action Generation – Army SBIR|STTR Program, https://armysbir.army.mil/topics/automated-course-action-generation/ 53. 30 Use Cases for Natural Language Generation (NLG) in the Defense Sector - Cogent Infotech, https://www.cogentinfo.com/resources/30-use-cases-for-natural-language-generation-nlg-in-the-defense-sector 54. Innovating Defense: Generative AI's Role in Military Evolution | Article - U.S. Army, https://www.army.mil/article/286707/innovating_defense_generative_ais_role_in_military_evolution 55. Reinforcement Learning-based Threat Assessment, https://arxiv.org/pdf/2503.02612 56. AI in Military Training and Simulation Market Revenue Trends and Growth Drivers | MarketsandMarkets, https://www.marketsandmarkets.com/Market-Reports/ai-in-military-training-simulation-market-5601755.html 57. Explaining Strategic Decisions in Multi-Agent Reinforcement Learning for Aerial Combat Tactics - arXiv, https://arxiv.org/pdf/2505.11311 58. Demystifying the Correlation of Forces Calculator, https://www.benning.army.mil/infantry/magazine/issues/2017/JAN-MAR/pdf/7)Spurlin_CoFCalculator_txt.pdf 59. Conversational AI over Military Scenarios Using Intent Detection and Response Generation, https://www.mdpi.com/2076-3417/12/5/2494 60. Tactical data link - Wikipedia, https://en.wikipedia.org/wiki/Tactical_data_link 61. Tactical Data Link (TDL): Enhancing Military Tactical Communications - Parraid, https://www.parraid.com/why-tdl-is-the-backbone-of-modern-military-tactical-communications/ 62. Link 16 - Wikipedia, https://en.wikipedia.org/wiki/Link_16 63. TACTICAL DATA LINKS, https://people.computing.clemson.edu/~jmarty/projects/lowLatencyNetworking/papers/TransportProtocols/TransportForTacticalDataNets/TacticalDataLinks-RAND.pdf 64. TACTICAL DATA LINK – FROM LINK 1 TO LINK 22 - ANMB, https://www.anmb.ro/buletinstiintific/buletine/2016_Issue2/MES/317-322.pdf 65. Tactical Communications for Ground-Based Air Defense (GBAD) - Bittium, https://www.bittium.com/wp-content/uploads/2025/05/Bittium-Tactical-Communications-for-Ground-Based-Air-Defense.pdf 66. Military Communication Systems: Securing Operations - Global sys, https://global-sys.com/news/military-communication-systems-for-modern-missions/ 67. 10 ways cybersecurity enhances military communication systems - Rocket.Chat, https://www.rocket.chat/blog/military-communication 68. The Integrated Tactical Network - Army University Press, https://www.armyupress.army.mil/Journals/Military-Review/English-Edition-Archives/May-June-2020/Blumberg-Int-Tactical-Network/ 69. Communications Challenges in DIL Environments | REDCOM, https://www.redcom.com/dil-environments/ 70. Maritime Operations in Denied, Intermittent, and Low ... - dodccrp.org, http://www.dodccrp.org/events/18th_iccrts_2013/post_conference/papers/088.pdf 71. Networking at the Tactical and Humanitarian Edge - Software Engineering Institute, https://www.sei.cmu.edu/blog/networking-at-the-tactical-and-humanitarian-edge/ 72. Army Tactical Network Quality of Service and Graceful Degradation Concept, https://cyberdefensereview.army.mil/CDR-Content/Articles/Article-View/Article/1134643/army-tactical-network-quality-of-service-and-graceful-degradation-concept/ 73. Robust communications for disconnected, intermittent, low-bandwidth (DIL) environments | Request PDF - ResearchGate, https://www.researchgate.net/publication/235257862_Robust_communications_for_disconnected_intermittent_low-bandwidth_DIL_environments 74. UFC 3-520-02 Facility Energy System Resilience and Reliability - Whole Building Design Guide, https://www.wbdg.org/FFC/DOD/UFC/ufc_3_520_02_2023_c1.pdf 75. FAULT-TOLERANT ARCHITECTURES FOR SPACE AND ..., https://www.cs.unc.edu/~anderson/teach/comp790/papers/Siewiorek_Fault_Tol.pdf 76. Using of Fault Tolerant Distributed Clusters in the Field of Command and Control Systems - DTIC, https://apps.dtic.mil/sti/tr/pdf/ADP010877.pdf 77. An Analysis of Redundancy Management Algorithms for Asynchronous Fault Tolerant Control Systems - NASA Technical Reports Server (NTRS), https://ntrs.nasa.gov/api/citations/19870019975/downloads/19870019975.pdf 78. Fault Tolerance Design and Redundancy Management Techniques. - DTIC, https://apps.dtic.mil/sti/tr/pdf/ADA090849.pdf 79. The Ethics of Artificial Intelligence in Defence - OII - University of Oxford, https://www.oii.ox.ac.uk/news-events/the-ethics-of-artificial-intelligence-in-defence/ 80. Navigating the AI battlefield: Opportunities and ethical frontiers - NRDC Italy, https://nrdc-ita.nato.int/newsroom/insights/navigating-the-ai-battlefield-opportunities--challenges--and-ethical-frontiers-in-modern-warfare 81. Transcending weapon systems: the ethical challenges of AI in ..., https://blogs.icrc.org/law-and-policy/2024/09/24/transcending-weapon-systems-the-ethical-challenges-of-ai-in-military-decision-support-systems/ 82. The comparative ethics of artificial-intelligence methods for military applications - PMC, https://pmc.ncbi.nlm.nih.gov/articles/PMC9510613/ 83. Human-in-the-loop (HITL) | Research Starters - EBSCO, https://www.ebsco.com/research-starters/computer-science/human-loop-hitl